

# Cracking Coding Interviews

## Longest Consecutive Sequence

**Mostafa S. Ibrahim**


*Teaching, Training and Coaching since more than a decade!*

*Artificial Intelligence & Computer Vision Researcher*

*PhD from Simon Fraser University - Canada*

*Bachelor / Msc from Cairo University - Egypt*

*Ex-(Software Engineer / ICPC World Finalist)*



# Leetcode [128](#) - Longest Consecutive Sequence

- Given an **unsorted** array of integers, return the length of the longest consecutive elements sequence.
- Input  $\Rightarrow$  Output
  - [100, 4, 200, 1, 3, 2]  $\Rightarrow$  4
    - Sequence [4, 1, 3, 2] when sorted [1, 2, 3, 4]  $\Rightarrow$  consecutive elements
  - [0, 3, 7, 2, 5, 8, 4, 6, 1]  $\Rightarrow$  9
    - The whole array when sorted is consecutive elements [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
- Signature
  - C++: `int longestConsecutive(vector<int>& nums)`
  - Python: `def longestConsecutive(self, nums: List[int]) -> int`
  - Java: `public int longestConsecutive(int[] nums)`

# Your turn

- Ask the right questions, if any, and state your assumptions
- Develop some test cases

# Questions & test cases

- The problem seems clear
  - A sequence => if sorted => consecutive items
  - If you feel confused, take your time to confirm what is being requested by the problem
- Several good cases are given already!

# Your turn

- Can we approach it with brute-force? If so, how?

# Your turn

- Can we approach it with brute-force? If so, how?
  - Not direct
- Hint: Given a number, can you use brute force to deduce the longest consecutive sequence from that number?
- E.g. for [100, 4, 200, 1, 3, 2] and value 2: can you get the answer, which is 3 [2, 3, 4]?
- Start with the initial value: keep checking if the increment to it exists in the array. Stop once you can no longer find it.
  - 2  $\Rightarrow$  there. Increment
  - 3  $\Rightarrow$  there. Increment
  - 4  $\Rightarrow$  there. Increment
  - 5  $\Rightarrow$  not there. Stop

# Your turn

- Now: for every number in the list compute its longest sequence and maximize
- How to check if a number is there?
  - Slow: loop to search in  $O(n)$ . Or, as we learned, we can do this quickly with a hash table

# Brute-force it!

```
int longestConsecutive(vector<int>& nums) {  
    unordered_set<int> st;  
    for (int i = 0; i < (int) nums.size(); ++i)  
        st.insert(nums[i]);  
  
    int ans = 0;  
    for (int val : st) { // modern C++ iterating  
        // The input has many duplicate values.  
        // This is iterating on all set elements (unique)  
  
        int len = 0;  
        while(st.count(val)) {  
            val += 1;  
            len += 1;  
        }  
  
        ans = max(ans, len);  
    }  
    return ans;  
}
```



# Brute-force it: Optimizations

- The last code was  $O(n^2)$ .
- We try every number
  - Compute the longest sequence
- Can we optimize more?
- We waste time trying every number in the sequence
- Assume input is [0, 1, 2, 3, 4, 5]
  - We get for 0 [0, 1, 2, 3, 4, 5]
  - We get for 1 [1, 2, 3, 4, 5]
  - We get for 2 [2, 3, 4, 5]
- But all are part of the SAME sequence. We need to process the whole sequence once!

# Brute-force it: Optimizations

- There are different ways to do that
- Here is an elegant way
- How can we recognize whether this number is the first number in the sequence?
- Always analyze the test cases
  - For [100, 4, 200, 1, 3, 2]
  - How to know in  $O(1)$  that 1 is the begin of a sequence but 2 is not?
  - If  $X$  is the first element in the sequence, then  $X-1$  won't be in the sequence
  - What is before 1? 0
  - Is 0 in the list? No. Then 1 is the first number in its sequence
  - What about 3? Before it is 2. Is 2 in the list? Yes. 3 is in the middle of a sequence!

# Optimized

- $O(n)$  time & memory

```
int longestConsecutive(vector<int>& nums) {  
    // Add to hashset to check if exists  
    unordered_set<int> st;  
    for (int i = 0; i < (int) nums.size(); ++i)  
        st.insert(nums[i]);  
  
    int ans = 0;  
    for (int val : st) { // modern C++ iterating  
        if(st.count(val-1))  
            continue; // NOT sequence first element  
  
        int len = 0;  
        while(st.count(val)) {  
            val += 1;  
            len += 1;  
        }  
        ans = max(ans, len);  
    }  
    return ans;  
}
```

# Another direction

- Another trivial solution is to sort numbers in  $O(n \log n)$ 
  - E.g. [1, 2, 3, 4, 100, 200, 201, 202, 500]
  - Now process them, and group consecutive numbers together
    - 1, 2, 3, 4
    - 100, 200
    - 200, 201, 202
    - 500
- This can't be improved to  $O(n)$ . Take care to avoid getting trapped yourself with a method that doesn't improve upon your previous solution
  - Consider other directions when you are stuck
  - Ask for a few hints if needed

# Interesting idea

- How can u formulate this problem as a graph problem?

# Interesting idea

- How can u formulate this problem as a graph problem?
- Each consecutive sequence is actually a chain of nodes
- So we can build the tree and then find the maximum connected component, e.g. using union-find
- How to build the graph Edges?
  - For number X
    - If  $X+1$  exists then we have edge  $(X, X+1)$
    - If  $X-1$  exists then we have edge  $(X-1, X)$
- Find solutions in your languages in the discussion tab. C++ [Code](#)
- *It is rare in interviews when you reinterpret the problem in another domain*

*“Acquire knowledge and impart it to the people.”*

*“Seek knowledge from the Cradle to the Grave.”*